

ASSIGNMENT 2 OF 2 · REQUIRED COVER SHEET · DUE SUNDAY 13 SEPTEMBER 2026, 23:59 SGT

A2 Team Self-Appraisal

Team ID: B-7 Section: B Problem chosen (A / B): A

Members: Goncalo Miranda · ZHENG YONGJIE · SUN YUCONG · JIN CHENG · NIU TONG · WANG HONGJUN

Repository URL: <https://github.com/gaamiranda/PE6201> Date: 20 September 2026

ONE SHEET PER TEAM, not one per member. Complete it together, in one sitting, after the work is done — it takes about twenty minutes. CHECK IT INTO YOUR A2 SUBMISSION FOLDER alongside the rest of the submission, by the A2 deadline of Sunday 13 September 2026, 23:59 SGT; do not email it. This sheet is required and is not graded, but a missing self-appraisal is treated as an incomplete submission.

1 · Rate your team against Rubric 1

Circle or mark one band per criterion. Be honest rather than generous — a team that marks itself Excellent everywhere tells the marker nothing, and a large gap between your rating and the marker's is the single most useful thing on this page.

Criterion	Weight	Excellent (4)	Proficient (3)	Developing (2)	Limited (1)
Conceptual Understanding	25%	✓ selected			
Technical Execution	30%		✓ selected		
Reasoning & Justification	25%	✓ selected			
Communication & Clarity	20%		✓ selected		

2 · The one decision you would defend hardest

Name a single design decision — a tool you cut, the autonomy setting you chose, the dependency rule behind your parallel calls, the model you settled on — and say what evidence made it. Four or five sentences.

Scoring the record separately from the decision, and letting a human judgement overrule a passing code check. It cost us an extra grading layer and twelve manual verdicts per model, and it is why our headline number looks worse than a decision-only rate would. The evidence is that the gap is real and large: on the scripted baseline 72 of 76 trials reached the expected decision but only 57 passed every check, and on mistral-medium-3 the same split is 72 against 53. A claims system that routes correctly and cannot say why has not done the job, and a decision-only rate would have hidden that on every model we tested. It also changed what we worked on: the failures cluster in naming the missing document, not in reading hostile narrative, which was the opposite of our prediction.

3 · The one you are least sure about

Every team has one. Naming it costs you nothing here and it is often the most interesting thing in the submission.

Shipping the v2 six-field descriptor rewrite. It costs 37.1% more input tokens and 35.3% more money per battery, and we cannot show it bought accuracy: on the same model and the same 76 trials, v1 scored 28 and v2 scored 31 — a three-trial gap against a standard error of about six, so roughly half a standard error and not distinguishable from noise. D7's three scripted arms scored identically, which is not independent evidence either, because replay returns saved responses. We kept it for inspectability — a size bound on returns, stated failure conditions, an explicit irreversibility flag — and that is a judgement about auditability, not a measured result. A paired replicated experiment would settle it and we did not run one.

4 · Your headline numbers

Copy these from your repository. If any cell is empty, that is worth noticing before you submit rather than after.

	Value		Value
Evaluation cases	42	Guardrail cases	10
Negative cases	17	Live runs per model (D4)	76
Pass rate (best model)	58/76 = 76.32%	Pass rate on negatives only	37/51 = 72.55%
Models in the battery	5 models, 6 batteries	Median turns per run	5

The live model YOU ran (D5b)	openai/gpt-4o-mini, v2	Your pass rate on it	31/76 = 40.79%
Turns saved by parallel calls	222 → 183 = 39 turns	Cost per successful task	US\$1.8014
Monthly cost at problem volume	US\$14,911.06	Break-even success rate	69.67%
Tokens per call — v1 → v2	265 → 540	Pass rate — v1 → v2, same model	28/76 → 31/76

Every member runs one live model, or the v1 pass — three models is the floor, N-1 for a team of N is the shape expected. Name YOURS. If you ran the v1 pass, write "v1 pass" and the model it sat on.

The last row is D2(b): the descriptor rewrite you measured. Two numbers each, before and after, on ONE model. An empty row here means the rewrite was never measured.

5 · Contribution

Who owned what. This should match CONTRIBUTIONS.md and the commit history. It is not used to scale marks — the team receives the team score — but it is one of the two signals the non-contribution clause relies on.

Member	Owned	Also contributed to
Goncalo Miranda	Agent loop, tool layer, vendor seam, autonomy gate (D1, D2(a), D2(c)); report §2	Coordination; live-path retry fix and battery-v2.1; price correction; D7 failure 1; gpt-4o-mini v2 battery
ZHENG YONGJIE	D0; report §1; assembly of all six sections	gpt-4o-mini v1 battery — the controlled prompt arm; fourteen overclaims found by reading D0 against the code
SUN YUCONG	Descriptors, the v1→v2 rewrite, the guardrail layer (D2(b), D3); report §5	deepseek-chat battery; found the missing live-path retry; injection results across models
JIN CHENG	Scripted backend and case loader (D4, D5(a)); report §6	gemini-2.5-flash-lite battery — the direct replay-versus-live comparison
NIU TONG	Evaluation harness and the D5(b) runbook (D4, D5(a)); report §3	llama-3.3-70b battery; the judgement workflow and its audit trail
WANG HONGJUN	Cost model, ledger and sensitivity analysis (D6); report §4	mistral-medium-3 battery; re-verified PRICES and caught a delisted model the day before the battery

Peer rating is separate from this sheet and comes AFTER it: NTU's Group Project peer-review function opens 13 September and closes Wednesday 16 September, 23:59 SGT. Each member completes it individually in NTULearn. It is a participation requirement.

6 · Declaration

Every member ticks every box. One sheet, signed by the team.

☒ Every member of this team can explain every block of code we submit — what it does, and why it is there. (Block, not line: this is about function, not syntax.)

☒ The pass rates, token counts, turn counts and costs we report are measurements we ran ourselves. None of them were supplied by an AI assistant.

☒ Our submitted harness runs end to end on the scripted backend, with no key and no network.

☒ We kept every fixture record we were given, unedited, and every case in our evaluation set carries a label. check_my_data.py passes on our data.

☒ No part of this submission is drawn from any member's End-of-Course Project, and no part of it will be submitted as part of one.

☒ This is our own work, and all sources, tools and assistance are attributed.

Signed on behalf of the team: _____

Date: 20 September 2026

How this is used

This sheet is not marked and cannot lower your grade. It is read alongside your submission for two reasons. First, how a team reads its own work is usually the fastest route to feedback that is actually useful. Second, where a team's self-rating and the marking diverge sharply, that gap gets an explicit comment — which is where the most useful feedback tends to sit.